

AVL Trees, Splay Trees, and Amortized Analysis



AVL Trees



Target : **Speed up searching (with insertion and deletion)**




Tool : **Binary search trees**




Problem : **Although $T_p = O(\text{height})$, but the height can be as bad as $O(N)$.**

〔Example〕 2 binary search trees obtained for the months of the year



Average search time = 3.5

Average search time of the skew tree = 6.5



Average search time = 3.1

A balanced tree

Adelson-Velskii-Landis (AVL) Trees (1962)




【Definition】 An empty binary tree is height balanced. If T is a nonempty binary tree with T_L and T_R as its left and right subtrees, then T is **height balanced** iff

(1) T_L and T_R are height balanced


(2) $|h_L - h_R| \leq 1$ where h_L and h_R are the heights of T_L and T_R , respectively.

The height of an empty tree is defined to be -1 .

【Definition】 The balance factor $BF(\text{node}) = h_L - h_R$. In an AVL tree, $BF(\text{node}) = -1, 0, \text{ or } 1$.




【Example】 Input the months Mar May Nov




👁 The trouble maker **Nov** is in the **right** subtree's **right** subtree of the trouble finder **Mar**. Hence it is called an **RR rotation**.

In general:

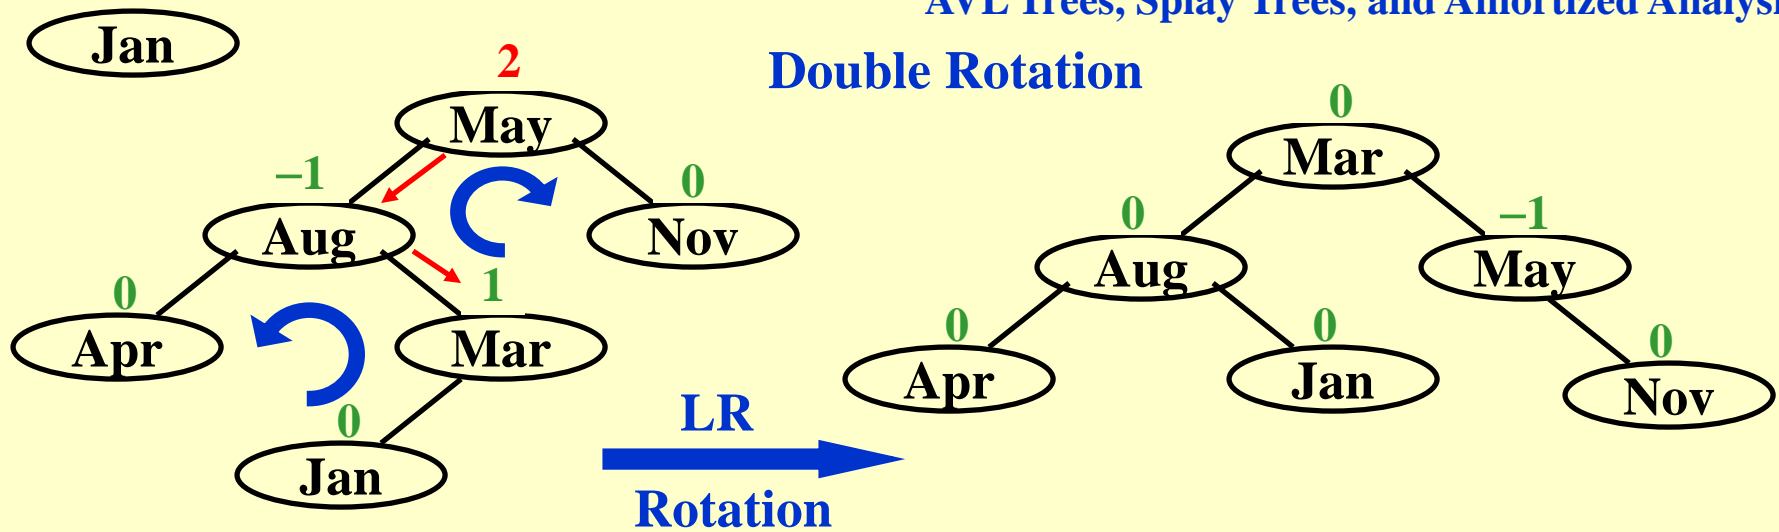


Aug Apr




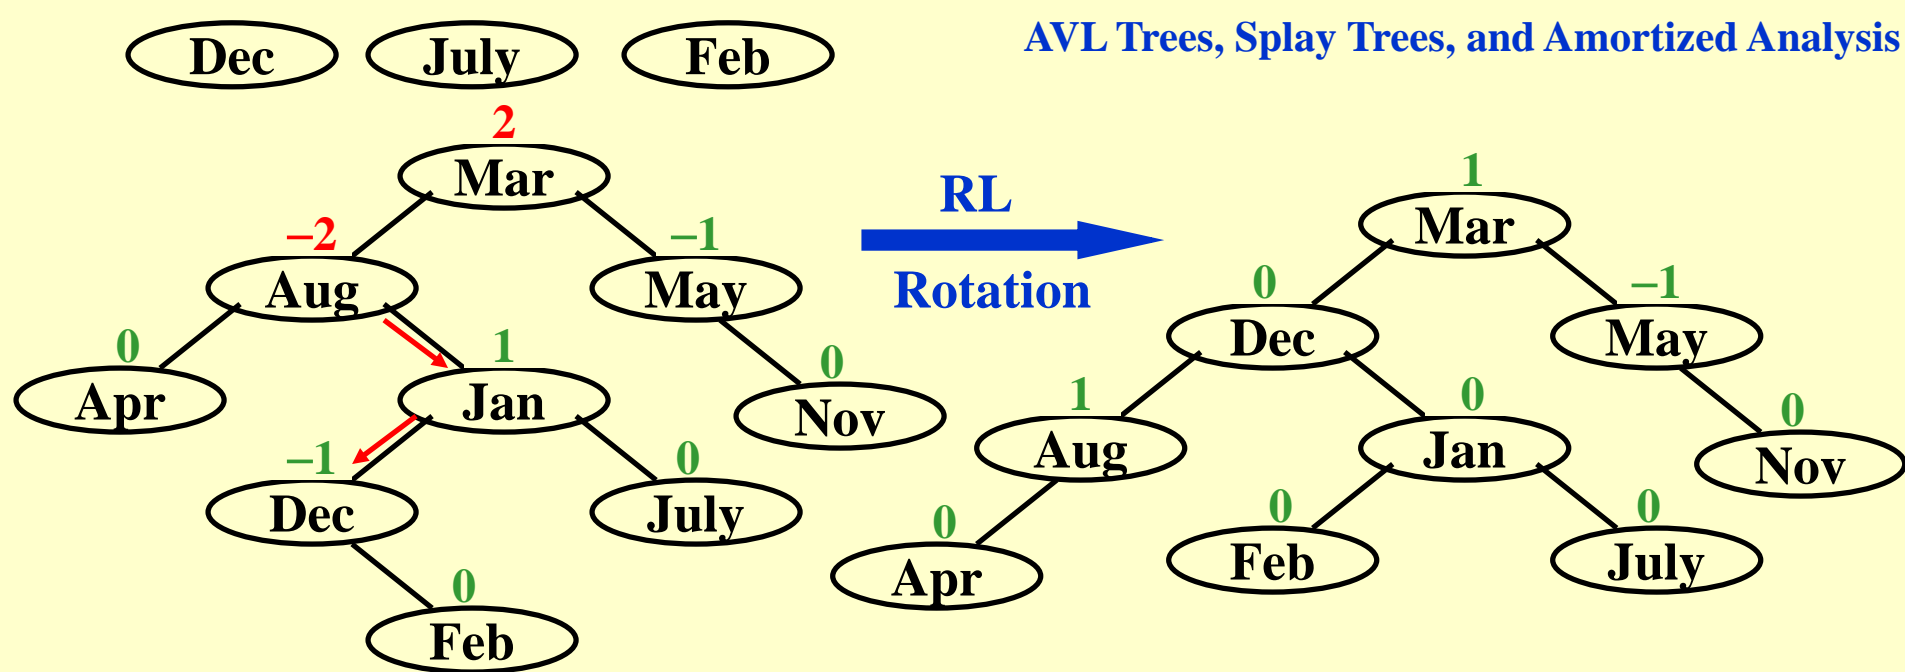
Discussion 1: What can we do now?

Double Rotation




In general:

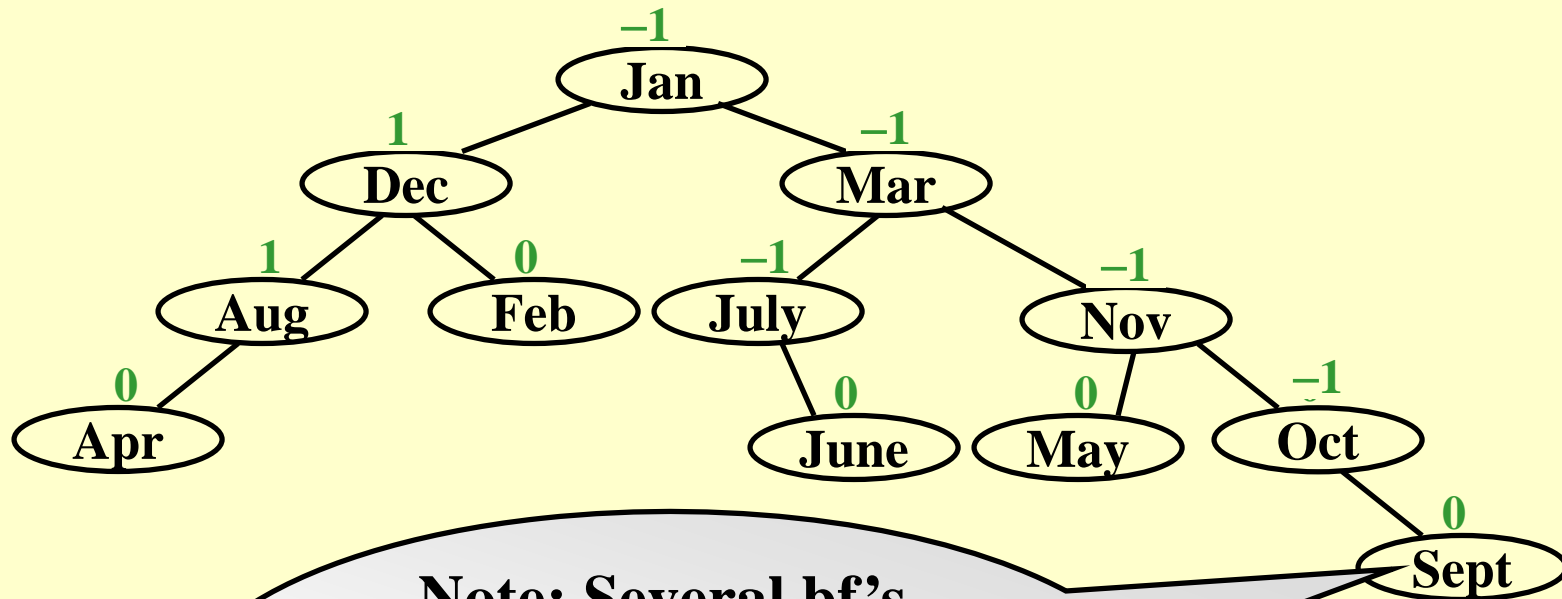




In general:



June Oct Sept




**Note: Several bf's
might be changed even if
we don't need to reconstruct
the tree.**

Another option is to keep a *height* field for each node.

Read the declaration and functions in [1] Figures 4.42 – 4.48

Let n_h be the minimum number of nodes in a height balanced tree of height h . Then the tree must look like



Fibonacci numbers:


$$F_0 = 0, F_1 = 1, F_i = F_{i-1} + F_{i-2} \text{ for } i > 1$$

$$\Rightarrow n_h = F_{h+2} - 1, \text{ for } h \geq 0$$

Fibonacci number theory gives that $F_i \approx \frac{1}{\sqrt{5}} \left(\frac{1 + \sqrt{5}}{2} \right)^i$

$$\Rightarrow n_h \approx \frac{1}{\sqrt{5}} \left(\frac{1 + \sqrt{5}}{2} \right)^{h+2} - 1 \Rightarrow h = O(\ln n)$$

Splay Trees




Target: Any M consecutive tree operations starting from an empty tree take at most $O(M \log N)$ time.

Sure we can – that only means that whenever a node is accessed, it must be **moved**.





One node takes $O(N)$ time
we can keep accessing it
times, can't we?



Idea: After a node is accessed, it is pushed to the root by a series of AVL tree rotations.



An even worse case:

Insert: 1, 2, 3, ... N **Find:** 1**Find:** 2..... **Find:** N 


$$T(N) = O(N^2)$$

Try again -- For any nonroot node X , denote its parent by P and grandparent by G :


Case 1: P is the root \rightarrow Rotate X and P

Case 2: P is not the root


Zig-zag




Double rotation




Zig-zig







Single rotation



Splaying not only moves the accessed node to the root, but also roughly halves the depth of most nodes on the path.



Insert: 1, 2, 3, 4, 5, 6, 7**Find: 1**

**Read the 32-node example
given in Figures 4.52 – 4.60**

Deletions:

☞ **Step 1:** Find X ;

X will be at the root.

☞ **Step 2:** Remove X ;


There will be two subtrees T_L and T_R .


☞ **Step 3:** FindMax (T_L) ;

The largest element will be the root of T_L , and *has no right child*.

☞ **Step 4:** Make T_R the right child of the root of T_L .

Are splay trees really better than AVL trees?






Research Project 1

Binary Search Trees (26)

This project requires you to implement operations on unbalanced binary search trees, AVL trees, and splay trees. You are to analyze and compare the performances of a sequence of insertions and deletions on these search tree structures.

Detailed requirements can be found at <https://pintia.cn/> .

Amortized Analysis



Target: Any M consecutive operations take at most $O(M \log N)$ time.

-- *Amortized* time bound

worst-case bound \geq amortized bound \geq average-case bound


Probability
is *not* involved

👉 Aggregate analysis

👉 Accounting method

👉 Potential method

👉 Aggregate analysis



Idea: Show that for all n , a sequence of n operations takes *worst-case* time $T(n)$ in the *worst case*, the average cost, per operation is

We can **pop** each object from the stack *at most once* for each time we have **pushed** it onto the stack

$$Total = O(n^2) ?$$

【Example】 Stack with **Pop**

```
Algorithm {
  while ( !IsEmpty(S) && k>0 ) {
    Pop(S);
    k - -;
  } /* end while-loop */
}
T = min ( sizeof(S), k )
```

Pop(int k, Stack S)

Consider a sequence of n **Push**, **Pop**, and **MultiPop** operations on an initially empty stack.

$$\text{sizeof}(S) \leq n$$

$$T_{\text{amortized}} = O(n)/n = O(1)$$

👉 Accounting method



Idea: When an operation's *amortized cost* \hat{c}_i exceeds its *actual cost* c_i , we assign the difference to specific objects in the data structure as *credit*. Credit can help *pay* for later operations whose amortized cost is less than their actual cost.

Note: For all sequences of n operations, we must have

$$T_{\text{amortized}} = \frac{\sum_{i=1}^n \hat{c}_i}{n} \geq \frac{\sum_{i=1}^n c_i}{n}$$

[[Example]] Stack with **MultiPop**(**int** k , Stack S)

c_i for **Push**: 1 ; **Pop**: 1 ; and **MultiPop**: $\min (\text{sizeof}(S), k)$

\hat{c}_i for **Push**: 2 ; **Pop**: 0 ; and **MultiPop**: 0

Starting with an empty stack Credits for

Push: +1 ; **Pop**: -1 ; and **MultiPop**: 0 for each +1

$\text{sizeof}(S) \geq 0$

The amortized costs of the operations may *differ* from each other

$$\Rightarrow O(n) = \sum_{i=1}^n \hat{c}_i \leq \sum_{i=1}^n c_i$$

$$\Rightarrow T_{\text{amortized}} = O(n)/n = O(1)$$

👉 Potential method



Idea: Take a closer look at the *credit* --

$$\hat{c}_i - c_i = \text{Credit}_i = \Phi(D_i) - \Phi(D_{i-1})$$

Potential function

$$\begin{aligned} \sum_{i=1}^n \hat{c}_i &= \sum_{i=1}^n (c_i + \Phi(D_i) - \Phi(D_{i-1})) \\ &= \left(\sum_{i=1}^n c_i \right) + \underbrace{\Phi(D_n) - \Phi(D_0)}_{\geq 0} \end{aligned}$$

In general, a good potential function should always assume its minimum at the start of the sequence.

[[Example]] Stack with **MultiPop**(**int** k , Stack S)

$D_i =$ the stack that results after the i -th operation

$\Phi(D_i) =$ the number of objects in the stack D_i

$$\Phi(D_i) \geq 0 = \Phi(D_0)$$

Push: $\Phi(D_i) - \Phi(D_{i-1}) = (\text{sizeof}(S) + 1) - \text{sizeof}(S) = 1$

$$\Rightarrow \hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) = 1 + 1 = 2$$

Pop: $\Phi(D_i) - \Phi(D_{i-1}) = (\text{sizeof}(S) - 1) - \text{sizeof}(S) = -1$

$$\Rightarrow \hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) = 1 - 1 = 0$$

MultiPop: $\Phi(D_i) - \Phi(D_{i-1}) = (\text{sizeof}(S) - k') - \text{sizeof}(S) = -k'$

$$\Rightarrow \hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) = k' - k' = 0$$

$$\sum_{i=1}^n \hat{c}_i = \sum_{i=1}^n O(1) = O(n) \geq \sum_{i=1}^n c_i \Rightarrow T_{\text{amortized}} = O(n)/n = O(1)$$

[[Example]] Splay Trees: $T_{amortized} = O(\log N)$


$D_i =$ the root of the resulting tree

$\Phi(D_i) =$ must increase by at most $O(\log N)$ over n steps, AND will also cancel out the number of rotations (zig:1; zig-zag:2; zig-zig:2).




$\Phi(T) = \sum_{i \in T} \log S(i)$ where $S(i)$ is the number of descendants of i (i included).

*Rank of the subtree
 \approx Height of the tree*

Why not simply use the heights
of the trees?



$$\Phi(T) = \sum_{i \in T} \text{Rank}(i)$$



【Theorem】 The amortized time to splay a tree with root T at node X is at most $3(R(T) - R(X)) + 1 = O(\log N)$.

Reference:

Data Structure and Algorithm Analysis in C (2nd Edition):
Ch.4, p.106-128; Ch.11, p.447-451; *M.A.Weiss 著、*
陈越改编, 人民邮电出版社, 2005

Introduction to Algorithms, 3rd Edition: Ch.17, p.
451-478; *Thomas H. Cormen, Charles E. Leiserson,*
Ronald L. Rivest and Clifford Stein. The MIT Press. 2009